


Carry a heavy hull block on one big transporter, or on several small ones coupled together?

Single-vehicle capacity or cooperative transport: shipyard transporter fleets sized at equal service and priced from tender awards


Single-vehicle capacity versus cooperative transport in the service-constrained sizing of shipyard block-transporter fleets

Chunrui Na, Qing Wang, Xiao Han, Jingzheng Yao, Kun Ma, Youngrong Kim
Harbin Engineering University · Nanyang Technological University
Target journal: Ocean Engineering (manuscript in preparation)

A · 1 × 500 t



B · 2 × 270 t



The whole study in one minute

1

The question

Coupled transporters can move blocks heavier than any single vehicle. Should a yard buy heavy units, light units that couple, or a mix? Existing studies schedule a given fleet.

2

What we did

Nine tiers and light-heavy mixes, each sized as the smallest count found to meet one service standard ($\geq 95\%$ on time, none over 120 min late) over 30 days in 36 conditions; priced from 33 tender awards.

3

The answer

Under shift labour the least-cost fleet couples at most one block in ten in **91.0%** of 1,620 decisions, and in all 1,350 where some tier carries every block. In 13 further conditions light units with matched heavy units win 95%.

4

Why


Coupled occupancy, alignment included, is a median 90% of a light fleet's extra transporter-time; partner waiting, 2%. Under shift labour a light tier can couple only about 2–18% of its workload.

Coupling is a remedy for exceptional blocks, not a purchasing strategy.

The order of this briefing (click to jump)

- | | | | | | |
|----------|--|-------|----------|---|-------|
| 1 | Background and question
How blocks move, what coupling is, what earlier work leaves open | p. 4 | 2 | Method
Equal service, scheduling and costing | p. 11 |
| 3 | Analytical model
Why single carry costs less per block; the coupling budget; four implications | p. 21 | 4 | Case and experiments
36 conditions, parameter sources, confirmation experiments | p. 25 |
| 5 | Results: what to buy
Count against capacity; the least-cost fleet | p. 30 | 6 | Why coupling costs
Occupancy, timing, the calibrated screen, capital against labour | p. 35 |
| 7 | Composition and service
Heavy units, lateness cap, on-time target, capacity loss | p. 39 | 8 | Validation
Exact benchmark, fresh days, search, parameters, a second yard | p. 47 |
| 9 | Discussion
When coupling pays, limits, five rules | p. 55 | | | |

Ships are built in blocks: blocks are made first and joined in the dock




A hull in the dock, welded together from many blocks (schematic)

- A large ship is divided into many **blocks**, built, outfitted and painted **in parallel** in different shops and joined in the dock
- A typical block is about 15 m × 15 m × 5 m and weighs 100–300 t; some exceed 500 t
- The price: blocks are **moved many times**, and a late move holds up the next process

This study is about those moves: which transporters, and how many, move them at least cost.


Each block is moved several times between shops, yards and the dock

A typical route of one block



- **Block masses** in published samples: 38–450 t, 204–399 t, 260–544 t
- **Transporters:** self-propelled hydraulic units; fleets in the literature are rated 200–550 t
- **Price:** about CNY 1.79–3.92 million per unit for 200–550 t (ex VAT, fitted to tender awards)
- **Crew:** four or five people per transporter; we use four
- **Workload:** on one recorded day four transporters moved 97 blocks in 16 h

How a transporter moves a block: drive under, lift, drive, set down



① Drive in empty

The block rests on supports; the vehicle drives into the gap underneath

② Lift hydraulically

The platform rises and lifts the block off the supports: this is loading

③ Drive loaded, slowly

It carries the block to the destination at about half its empty speed

④ Set down, leave

The platform lowers the block onto new supports (unloading), and the empty vehicle heads for the next block

Rated capacity

Sold by tier; one unit cannot carry a block above its rating

Speed

270 t model: 12 km/h empty, 6 km/h loaded


Loading / unloading

10 min each; a mass-dependent formula also exists; loading in practice often takes 30–60 min

Crew

Four or five per vehicle (driver and guides)


Coupling: two or three units side by side drive as one large vehicle



- **Capacities add:** two 270 t units lift up to 540 t, three up to 810 t; the set runs under synchronised control
- **All must be there:** the first to arrive waits for the last
- **Alignment takes time:** no measured value is published; we assume 10 min and test 0–40 min
- **Held throughout:** from loading to unloading, each unit and its crew can do nothing else
- At most three units per team (assumed; two also tested)

Coupling lets light units carry heavy blocks, but one block then takes two or three times the vehicles and crews: cheaper capital, more vehicle-time.

The problem in one figure: block logistics, three fleet options, one comparison



- ① **Block logistics:** each block moves several times between shops, stockyard and building dock
- ② **Three fleet options:** A heavy units only; B light units coupled for heavy blocks, with alignment δ before each coupled lift; C light units with a few heavy units, and the question of how many
- ③ **The comparison:** the same service $\rightarrow$ the smallest count found by scheduling $\rightarrow$ the least daily cost; a fluid occupancy model explains the cost differences

Every fleet is held to the same service before costs are compared.

Earlier work schedules coupled lifts and compares vehicle counts, but never weighs coupling against buying a larger carrier

Shipyard block logistics

- Less empty travel and interference Roh & Cha 2011
- Heterogeneous transporters, delivery limits Kim & Joo 2012; Joo & Kim 2014
- Workload balance, set-ups, precedence, routing Park & Seo 2012
- Paths: distance and number of turns Wang et al. 2020
- Energy and time windows, barge loading, stockyard sequencing Guo et al. 2024; Kweon et al. 2026; Wang & Yin 2026
- More units above 320 t, from fleet utilisation Liu et al. 2022

Cooperative transport

- Vehicles synchronised for one large cargo Hu & Wei 2018
- Multi-vehicle one-cargo for 200–500 t blocks; time and CO₂ Jiang et al. 2021
- Two transporters share a block in digital-twin scheduling Sun et al. 2024
- One job-shop transport needing several robots Gayon et al. 2026
- Operation, movement and load synchronisation Drexl 2012; Soares et al. 2024

Vehicle numbers and fleet mix

- Shipyard vehicle counts compared; one type, divisible tasks Zhou et al. 2025
- Types and numbers with routes; fixed costs Golden et al. 1984; Hoff et al. 2010; Koç et al. 2016
- AGV and mobile-robot fleets Vis et al. 2005; Vis 2006; Fragapane et al. 2021; Aziez et al. 2022
- Synchronised visits under uncertainty Hashemi Doulabi et al. 2020
- Robot-team composition from simulated schedules Pan & Yu 2024

Analogues, service and search

- Couplable warehouse robots: identical units, only how many Albert et al. 2025; Chaikovskaia et al. 2026
- Modular transit coupling along a corridor Cao et al. 2022
- Fleet size at a stated service level Eskandari & Mahmoodi 2016
- Production scheduling with conflict-free AGV routes Nishi et al. 2011
- ALNS operator performance depends on the problem Voigt 2025

The gap: shipyard work schedules coupled loads and compares vehicle counts; routing and robot-team work link resource choice to performance. Still open: buying single-carrier tonnage or coupling, for indivisible blocks at the same service. We keep teams, alignment, empty travel and due dates in daily scheduling, price fleets from tenders and crew labour, and test how far decisions rest on search and sampled days.

What this study asks, and what makes it hard

Q1 What to buy §5.1

At equal service, which fleet is cheapest, and how much does it couple?

Q2 Why §3.3, §5.2

What makes coupling expensive, and how much coupling can a light tier afford?

Q3 Composition and service §5.3, §5.4

How many heavy units? How should service be specified? Is coupling insurance against breakdowns?

Q4 How firm §5.5

How far do search error, sampled days and parameters move the decision?

Four complications

- **Service:** a fleet can meet an on-time target by letting its hardest blocks slip, so the answer depends on how service is specified
- **Mixes:** light units with a few heavy units offer a middle way between the extremes
- **Insurance:** coupling is often argued for as cover for a broken heavy unit
- **Search error:** counts come from a heuristic, and a rival that qualifies with one transporter fewer can overturn a cost ranking


The core idea: hold every candidate fleet to the same service, then compare costs

An analogy

You cannot compare ten buses running 80% on time with twenty minibuses running 99% on time by price alone. First give each fleet the same punctuality, then see which is cheaper.

- For each fleet type, scheduling finds the **smallest count that meets the service, K^*** , an upper bound on the true minimum
- Then cost K^* units: capital plus crew labour per day

$$\geq 95\% \text{ on time} \cdot T_{\max} = 120 \text{ min}$$





- **95% on time:** at least 2,765 of the 2,910 blocks of 30 days on time (at most 145 late)
- **$T_{\max} = 120$ min:** no block may finish more than 120 min late, one release-batch interval, so a late block cannot overrun the next batch

Why not just divide the total work by one vehicle's hours?

A small example: four blocks ready at 8:00, all due by 9:00, each move takes 1 h

Totals only: $4 \text{ vehicle-h} \div 16 \text{ h} = 0.25 \rightarrow$ one vehicle? By due time: four vehicles at once




- **Work arrives in batches** with their own due times, so demand is uneven
- **Vehicles drive empty between tasks**, and the order changes that distance
- **A coupled lift needs several free vehicles at once**; one late arrival holds the team

So each day's **actual timetable** must be built. The workload-only count K_p is within one transporter of the scheduled count for 319 of 384 sized fleets, but under tight due dates fleets that couple more than 60% of blocks need 10.2% of K^* more.

One working day: 97 blocks released in eight batches, each with a due time

One generated day (uniform masses · short handling · empirical due dates)



- **97 blocks in 16 h:** the size of a recorded working day (8:00 to midnight)
- **Each line is a block:** left end release, right end due time; height is mass
- **Due times:** 150, 180 or 210 min after release, with probabilities 0.6 / 0.2 / 0.2
- **Release:** a batch every 2 h, about 12 blocks (assumed)
- **Orange** blocks exceed 270 t: a 270 t fleet must couple all of them (62 that day)

Each condition has 30 such days, and every fleet faces exactly the same 30 days.

Who may carry together? Only minimal feasible teams

The team's ratings must add up to the block's mass, and removing any member must make it fall short. Extra units would only waste vehicles and crews.

$$\sum Q_k \geq m_i \quad \cdot \quad \sum Q_k - \min Q_k < m_i \quad \cdot \quad |S| \leq 3$$

- A block one unit could carry may still go to a light team (flexible coupling, the main rule); the overweight-only rule is compared separately
- For a homogeneous fleet the two rules coincide; only light-plus-heavy mixes differ
- Every minimal team is assumed physically feasible (control, load sharing and road width not modelled), which favours coupling

Try it

Transporter tier (rated capacity)


200 t 250 t **270 t** 300 t 325 t 380 t 425 t 500 t 550 t

Block mass  450 t



2 coupled: $2 \times 270 = 540 \text{ t} \geq 450 \text{ t}$; one fewer gives only $270 \text{ t} < 450 \text{ t}$. Holds 2 vehicles · 8 crew, plus 10 min coupling.


The price of coupling: every member must be there, and all are held for the whole move



One move up close

- **Alone:** the vehicle loads on arrival and leaves right after unloading
- Example: a 420 t block takes 43 transporter-minutes on one 425 t unit, but $2 \times 55 = 110$ on two coupled 270 t units
- **Coupled:** the first unit waits for the second (white), then both align (yellow)
- A light block waiting for one of them finishes at 98 min instead of 43

The same day and the same 97 blocks for two fleets: their actual timetables



	500 t × 5	270 t × 9
Vehicles K*	5	9
On time that day	97 / 97	96 / 97
Blocks coupled	0	62
Vehicle work (vehicle-h)	70.0	145.0
· empty travel	5.7	16.7
· waiting for team	0.0	1.9
Worst delay (min)	0	26
Fleet price P (270 t = 1)	8.15	9.00
Daily cost (labour-h, $\theta = 12.4$)	421	688

The 270 t fleet needs nine units and twice the vehicle-hours for the same work, coupling 62 blocks; in this condition even its capital is higher.

Each row is a vehicle and each bar a task; hover to see the block. Both are stored qualifying schedules at each fleet's smallest count, re-simulated to check.


Sources: schedules of this study (uniform masses · short handling · empirical due dates, first day).

How timetables are found: far too many to try, so we search cleverly

97! $\approx 10^{152}$

The order of 97 blocks alone has about 10^{152} possibilities, before choosing any teams.

So we use **adaptive large neighbourhood search (ALNS)**, improving a workable timetable step by step.



⌚ 786 rounds per run (38 s on average), exactly repeatable from its seed; the start is the best of 5,000 random chronological constructions

How two timetables are ranked

① blocks over 120 min late → ② late blocks → ③ total lateness → ④ crew time. One more block on an earlier item outweighs any gain on later ones. A team's members are updated together at every step, so coupled lifts stay synchronised.

Capital and labour on one scale: labour-hours

Vehicle

CNY 2.21 M

Price of one 270 t transporter (ex VAT, fitted to award prices)



CNY 329 k / yr

Spread over 8 years at 4% interest (capital recovery factor 0.1485)



1,316 CNY / day

250 working days a year

Labour

79,182 CNY / yr

Average wage of manufacturing production workers in China (NBS, 2025)



× 1.3 ÷ 2,000 h

Plus 30% employer on-costs (social insurance etc.), 2,000 hours a year



51.5 CNY / labour-h

Cost of one worker for one hour

Price ratio θ

$\theta = 1,316 \div 51.5 = 25.6$

One 270 t transporter's daily capital cost = 25.6 labour-hours

Compare

64

One vehicle's crew per day: 4 workers × 16 h = 64 labour-hours

Total cost

$C = \theta \cdot P + L$

P: the fleet's price in units of one 270 t transporter; L: the fleet's labour-hours per day

With Shanghai manufacturing wages (CNY 120.1 per labour-hour) $\theta = 11.0$. Over the combinations of life, interest, working days, on-cost and the two wages, θ lies in 5.1–30, and every result is reported over that range.

Sources: 270 t price fitted to 33 public tender awards (ex VAT); wages National Bureau of Statistics of China, release of 15 May 2026, 2025 average wages (large enterprises · manufacturing · production workers, CNY 79,182), and China Statistical Yearbook 2025, Table 4-12 (Shanghai manufacturing, CNY 184,833); on-cost from official contribution rates totalling ≤ 39.4% (Shanghai HRSS 2025; State Council 2019 No. 13; MOHRSS 2015 No. 71; MOHURD 2016 No. 74); crew of four Park & Seo (2012). Service life, interest, working days and 2,000 h a year are our decisions; θ is also tested over 0.5–5,000 (paper §3.2, Table 3).

Four ways to count crew labour; 180 cost scenarios per condition, 6,480 decisions in all

Shift labour

$$64 \cdot K$$

Every transporter staffed by a crew of four for the 16-h day. The primary measure: it is the staffing a yard must provide, and it grows with K.

Shift + overtime

$$64 \cdot K + \text{after-shift}$$

Adds 1.5×4 crew-hours for every transporter-hour worked after the shift.

Crew-vehicle hours

$$4 \times \text{vehicle-hours}$$

Four crew for the hours each transporter actually works (empty travel, waiting and service).

Crew-team hours

$$4 \text{ per team when coupled}$$

Four per vehicle for empty travel and waiting, but only four per team during a coupled lift, as if one crew could drive the set: the most favourable to coupling.


$$5 \text{ price curves} \times 9 \text{ values of } \theta \times 4 \text{ labour measures} = 180 \text{ cost scenarios per condition} \cdot \times 36 = 6,480 \text{ decisions}$$

- In each scenario the fleet with the lowest C_X is the **least-cost fleet**; the runner-up's excess over it is the **margin**
- The last three measures re-price **the same schedules**; schedules are not re-optimised for each

The 6,480 decisions re-price the same 384 sized fleets under different cost parameters; they are **not independent experiments**, and shares of decisions are descriptive frequencies over the grid, not probabilities.

Prices: a vehicle-day of capital costs less than a crew-day

Unit prices of 39 transporters (33 public tender awards, 2022–2026, ex VAT)




The 27 records within 200–550 t give $\alpha = 0.84$ (95% CI 0.72–0.96); over all 33 records a lack-of-fit diagnostic ($p = 0.039$) argues against a single power law, so five price curves are used: affine (main), piecewise and $\alpha = 0.72, 0.84, 0.96$.

Per transporter-day: capital (bars) against one crew (dashed), in labour-hours

Price ratio θ

 30.0



At $\theta = 30.0$, the dearest tier, 550 t, costs 53 labour-hours of capital a day, still below one crew's 64. Saving a transporter beats moving up a tier.

Single-carry dominance: for every block, one large-enough transporter costs less than any coupled team

One 420 t block (affine prices, $\theta = 12.4$, shift labour, short handling)

	ONE 425 T ALONE	TWO 270 T COUPLED
Capital π	1.43	1 + 1 = 2.00
Crews	1	2
Occupancy	43 min	2 × (43 + 10) min
Cost per block	3.7 labour-h	8.4 labour-h

cost per block = $(w_i + \delta \cdot 1[|S| > 1]) \div H \times \sum_{k \in S} \gamma(Q_k),$
 $\gamma(Q) = \theta \cdot \pi(Q) + \ell$

Fluid approximation: transporters divisible and fully usable over the day, counting only the vehicle-time each block takes; due dates, integer counts and schedule geometry are ignored.

Proposition 1: with sub-additive prices ($\pi(a + b) \leq \pi(a) + \pi(b)$); all five price curves over 200–1,650 t), if the menu has a tier Q^* with $m_i \leq Q^* \leq$ the team's capacity, one transporter of Q^* carries block i at a strictly lower fluid cost than any minimal team of $g \geq 2$.

- **No more capital:** one larger unit costs no more than the members together
- **$g - 1$ fewer crews:** crews are paid per vehicle
- **No alignment:** every member is held δ longer

It holds with labour alone ($\theta = 0$) and with capital alone ($\ell = 0, \delta > 0$); even at $\delta = 0$ a coupled lift adds $(|S| - 1) \cdot w + |S| \cdot \delta$ of transporter-time.


Corollary 1: if every block has a just-large-enough tier, the fluid-optimal fleet couples nothing and is mass-matched.

The coupling budget: how much of its workload can a light tier couple and still beat the covering tier?

Price ratio θ  12.4

Alignment δ (min)  10.0

Occupancy $\bar{w}$ (min)  43



 κ : how much dearer the larger unit is per day

 h^* : the largest coupled share of workload the light tier can afford

$h^* = \kappa \div (1 + 2\delta / \bar{w}) = \kappa \div 1.47$; $\delta = 10.0$ min, $\bar{w} = 43$ min (covering fleet's mean occupancy per block: about 43 with short, 63–85 with mass-dependent handling, about 80 with long loading); $\alpha = 0.84$ power law, shift labour

$$h \cdot (1 + 2\delta / \bar{w}) < \kappa \Leftrightarrow h < h^* = \kappa / (1 + 2\delta / \bar{w})$$

- h : the coupled share of single-carrier occupancy (workload, not count); $\kappa = \gamma(Q_{hi}) / \gamma(Q_{lo}) - 1$, how much dearer the covering tier is per day
- Labour dilutes the capacity premium: for $\theta = 5.1$ –30, κ is only 3.4–14.8% (425 vs 270 t) and 5.3–22.5% (550 vs 300 t)
- Covering fleets occupy a transporter about 43 min per block with short handling, 63–85 with mass-dependent handling and about 80 with long loading; with $\delta = 10$ min the factor is 1.47 or 1.24–1.32
- So $h^* \approx 2$ –18% under shift labour and the fitted power law (1.9–21.8% across the five price curves); yet in the 36 conditions the 200–380 t tiers couple **5–100%** of blocks

The budget is small because labour dilutes the capacity premium.

What the model leaves out, and the four implications the results test

In real operations but not in the fluid model (each can favour coupling)

Menu gap

If no tier can carry a block, coupling is necessary.

Integrity

Transporters are indivisible: a heavy tier kept for a thin heavy tail may stand idle; a light fleet with one or two heavy units is the integer counterpart of a mass-matched fleet.

Timing


Release peaks and due dates need slack beyond average workload, and a coupled lift needs several vehicles free at once.

Four implications

- **1** Where some tier carries every block, least-cost fleets couple little
→ Results: what to buy
- **2** A light tier beats the covering tier only when its coupled occupancy is below h^*
→ Why coupling costs: the calibrated screen
- **3** The heavy units of the cheapest light-heavy mix track the heavy blocks' share of occupancy
→ Composition: the heavy-share rule
- **4** Coupling-reliant fleets can win only where the budget does not bind or is large: a menu gap, near-zero alignment, or capital dear relative to labour
→ Capital against labour; when coupling pays

Two kinds of heavy: blocks above a light tier but not the largest can be served by buying capacity and govern the purchase; blocks above every tier must be coupled whatever the fleet and only set a floor on coupling.

Evaluation procedure: specify service, schedule each fleet, scan its count, price it at K^*



- ① **Service specification**, the same for every fleet: $\geq 95\%$ on time pooled over 30 days, and every block within $T_{\max}$
- ② **Schedule**: ALNS with single and coupled lifts; rank by over $T_{\max} > \text{late} > \text{lateness} > \text{crew time}$
- ③ **Count scan**: the smallest K found to qualify; three searches a day at $K^* - 1$; a schedule for K stays feasible for $K + 1$
- ④ **Price**: $C_X = \theta \cdot P_X + L_X$; the least-cost fleet in each cost scenario

Search error enters at step ③, where K^* is an upper bound, and can change the decision at step ④; the validation part tests both. All fleets face the same blocks, locations and start points (common random numbers); a fleet of K uses the first K of 40 start points.

Case: 36 conditions on one digitised yard, plus a second yard rebuilt from published data

Eight flows and their relative weights (assumed)

Assembly shop → Outfitting shop	0.5
Assembly shop → Paint shop	0.5
Outfitting shop → Paint shop	0.5
Paint shop → Block yard	0.7
Paint shop → Assembly yard	0.3
Block yard → Assembly yard	0.7
Assembly yard → Dock	0.2
Assembly shop → Assembly yard	0.05

The road network is the digitised layout of a large shipyard (not released); travel times are static shortest paths at 12 km/h empty and 6 km/h loaded (270 t data sheet), without traffic interference; pickup and drop-off points are drawn within the zones; eight batches a day, one every 2 h; each transporter starts at a random shop entrance.

36 × 30

6 mass profiles × 3 handling regimes × 2 due-date regimes; 30 days each: 1,080 days, 104,760 block moves

11

reference fleet types: 9 tiers (200–550 t) and 270 t units with one or two heavy units (scheduled as 550 t, priced at the smallest tier that carries every block they carry alone)

384 · 6,480 sized fleets (200 and 250 t not run for the heavy tail) · decisions

214,210

scheduling runs (51,480 in the main experiment), about 2,570 core-hours; each seeded and budgeted by iterations, not time, so it repeats exactly

50

second yard: the published week of transfers, distance table and five units (250–420 t) of a Shanghai cruise-ship yard


Sources: 97 blocks a day Roh & Cha (2011); speeds Suzhou Dafang DCY270 data sheet; flow categories and weights are assumed; second-yard data Liu et al. (2022), Tables 4–5 and supplementary data S1; run counts are ours (paper §4.1, §4.3).

36 conditions = 6 mass profiles × 3 handling regimes × 2 due-date regimes

Six mass profiles: masses of the 2,910 blocks of 30 days (orange: above 270 t)


Jiang • 38–450 t

68 masses from the three instances of Jiang et al. (2021), resampled




Light-skewed • 100–500 t

90% at 100–300 t (Park & Seo 2012), 10% at 300–500 t (assumed)




Uniform • 100–500 t

Uniform 100–500 t (the rule of the public instances of Kweon et al. 2026)




Heavy tail • 100–800 t

Light-skewed plus a 10% tail at 500–800 t: some blocks exceed every tier (assumed, a stress test)




Roh–Cha • 260–544 t

The ten block masses of Roh & Cha (2011) Table 2, resampled



Liu • 204–399 t

The 50 block masses of one week in Liu et al. (2022), resampled



Three handling regimes

- **Short:** 10 min each (Liu et al., 2022)
- **Mass-dependent:** $U(10, 20) + 0.05 \times \text{mass} - 5$ min each (Kim & Joo, 2012)
- **Long loading:** 30–60 min loading, 10 min unloading (Liu et al., 2022)

Two due-date regimes

- **Empirical:** 150 / 180 / 210 min with probabilities 0.6 / 0.2 / 0.2 (Roh & Cha, 2011)
- **Tight:** 120–150 min (assumed, tighter than the shortest published 150 min)

Parameter sources (1): values taken from the literature or public data

PARAMETER	VALUE	SOURCE (PAGE / TABLE)
Moves per day · shift	97 · 16 h	Roh & Cha (2011), p. 3246
Baseline due times	150 / 180 / 210 min (0.6 / 0.2 / 0.2)	Roh & Cha (2011) Table 2
Mass samples	38–450 t · 260–544 t · 204–399 t	Jiang et al. (2021) Tables 5/A1/A2; Roh & Cha (2011) Table 2; Liu et al. (2022) supplementary data
Uniform masses · footprint	100–500 t · 5–15 m by 10–30 m	Kweon et al. (2026) public instances
Handling	10 min each · mass-dependent · 30–60 min loading	Liu et al. (2022) §4.1.1, §4.2.2; Kim & Joo (2012) §4
Speed	12 empty · 6 loaded km/h; by model 10/5, 12/5	Suzhou Dafang DCY270, DCY380, DCY500 data sheets
Tiers	200–550 t	Jiang et al. (2021) Table 5; Roh & Cha (2011) Table 1; Liu et al. (2022) Table 4
Crew per vehicle	four (published: four or five)	Park & Seo (2012), p. 392
Prices	39 units, CNY 1.34–4.42 M, 90–600 t	33 public tender awards, 2022–2026 (our Supplementary S1)
Wages	CNY 79,182 a year (national) · 184,833 (Shanghai)	National Bureau of Statistics (2025); China Statistical Yearbook 2025, Table 4-12
On-cost	30–40%	official contribution rates, ≤ 39.4% in total

Sources: each entry was checked against the original page and sentence; full references on the last slide.

Parameter sources (2): values we had to set, and how each was tested

PARAMETER	SETTING	TESTED AT
Release	eight batches, 2 h apart	identical for all fleets
Tight due dates	U(120, 150) min	a condition beside the published due dates
Heavy part of light-skewed · heavy tail	10% × 300–500 t · 10% × 500–800 t	stress conditions; the heavy tail is analysed separately
Alignment δ	10 min	0, 2.5, 5, 20, 30, 40 min: single-carry share 88.6% at $\delta = 0$, 100% at 20–40 min
Team size · turns	$\leq 3 \cdot 0$ s	at most two; 30 s per turn of 60° or more; 30 or 60 s per turn for coupled teams
One speed for all tiers	12 / 6 km/h	manufacturer speeds 10/5, 12/5 km/h; speed × 0.75, × 0.5; Liu's 50/30 m/min
Handling multiple	× 1	× 0.5, 0.67, 1.5, 2; plus ±50% random noise
Service	95% · 120 min	caps of 30, 60, 240, 480 min and none; on-time targets of 90% and 98%
Capital recovery · hours	8–15 y, 4–8%, 250–300 days; 2,000 h a year	θ reported over all of 5.1–30, then 0.5–5,000
Sampled days	30 per condition	2,000 bootstrap resamples; 30 fresh days
Search budget	786 rounds; 5,000 constructions	counts after constructions, 100 and 300 rounds; six more searches a day at $K^* - 1$; exact counts on small instances

Sources: results of these tests are in the composition-and-service and validation parts (paper §5.3–5.5, Supplementary S2–S5); manufacturer speeds Suzhou Dafang DCY270, DCY380, DCY500 data sheets; 50/30 m/min Liu et al. (2022) §4.1.1.

Confirmation experiments: the model's implications tested on new instances, and an exact benchmark

Fresh days

30 days never used elsewhere; the 42 fleets contending for least cost in seven conditions

Extended mixes

13 further conditions: 300 t units with $K_H^* - 2$ to $K_H^* + 2$ covering heavy units (500, 550 or 425 t)

Slow speed and long handling

425 t, 500 t and mixes added at slow speeds and at handling $\times 1.5$ and $\times 2$; unresolved fleets scanned up to 40

Capacity loss

one fewer transporter in each capacity class, start positions reassigned; redundancy estimated from qualifying counts

On-time target

the 42 fleets re-sized at 90% and 98%

Exact benchmark


80 single-batch slices of 12–13 blocks from ten days in eight conditions; CP-SAT computes the fewest late blocks at each count; 70 sized fleets

$$\text{Heavy-share rule: } K_H^* = \lceil K_{\text{cov}} \cdot h_H \rceil$$

K_{cov} : the count of the covering homogeneous tier; h_H : the share of loading and unloading time taken by blocks above 300 t. Three further analyses use the stored runs: a redundancy estimate from the qualifying counts of reduced compositions, schedule re-selection for each labour measure, and a block-by-block decomposition of transporter-time.

Question 1: larger tiers need fewer transporters until one tier carries every block alone; beyond it, no further saving

Smallest qualifying count K^* (empirical due dates); dashed: the covering tier, the smallest tier that carries every block alone



- **Liu masses**, short handling: 12 units of 200 t, 9 of 270 t, 6 of 380 t, 5 of 425 t or more
- **Roh-Cha masses**: 15, 12, 8 and 5 units of 200, 270, 425 and 550 t
- Slower handling raises every line (up to 26 units of 200 t under Roh-Cha masses, 30 with tight due dates); the shape stays
- Beyond the covering tier extra capacity buys nothing: under Liu masses 425, 500 and 550 t need the same count in all six conditions
- Below it, the masses fix the coupled share: the 200–380 t tiers couple 5–100% of blocks

A light fleet's extra units are the price of coupling.

Question 1: the least-cost fleet carries almost every block alone

91.0%

of 1,620 shift-labour decisions: the least-cost fleet couples $\leq 1/10$ (all 1,350 where a covering tier exists)

57.5%

of all 6,480 decisions: the winner couples no block

24.1%

won by light units with one or two heavy units, coupling $\leq 1/10$: each carries its own blocks

1.2%

won by a light tier (≤ 325 t), only in the heavy tail or the hardest Jiang condition

By labour measure (share coupling $\leq 1/10$, 1,620 decisions each)



* one crew drives a coupled set, the measure most favourable to coupling; outside the heavy tail, all 32 decisions won by a coupling-reliant fleet use it.

- Over all four measures 90.8% (5,882 of 6,480); homogeneous fleets carrying alone win 66.7%
- The threshold is not critical:** $\leq 2.5\%$ in 74.5%, $\leq 15\%$ in 97.3%; price curves 90.0–91.6%; θ 5.1 $\rightarrow$ 30: 93.1% $\rightarrow$ 87.6%

Every minimal team is assumed physically feasible, which favours coupling, so that assumption does not explain these shares.

All 36 conditions × 9 price ratios: which fleet is cheapest

		5.1	6.4	7.9	9.9	12.4	15.4	19.3	24.0	30.0
Jiang	short · empirical	6	6	6	6	6	6	6	6	6
	short · tight	0	0	0	0	0	0	0	0	0
	mass-dep. · empirical	1	1	1	1	1	1	1	1	1
	mass-dep. · tight	1	1	1	1	1	1	1	1	1
	long · empirical	1	1	1	1	1	1	1	1	1
	long · tight	0	0	0	0	0	0	0	0	0
Light-skewed	short · empirical	3	3	3	3	3	3	3	3	3
	short · tight	4	4	4	4	4	4	4	4	4
	mass-dep. · empirical	2	2	2	2	2	2	2	2	2
	mass-dep. · tight	0	0	0	0	0	0	0	0	3
	long · empirical	0	0	0	0	0	0	0	2	2
	long · tight	0	0	0	0	0	5	5	5	5
Uniform	short · empirical	0	0	0	0	0	0	0	0	0
	short · tight	0	0	0	0	0	0	0	0	0
	mass-dep. · empirical	0	0	0	0	0	0	0	0	0
	mass-dep. · tight	0	0	0	0	0	0	0	0	0
	long · empirical	0	0	0	0	0	0	0	0	0
	long · tight	0	0	0	0	0	0	0	0	0
Heavy tail	short · empirical	9	9	9	9	9	9	9	9	9
	short · tight	10	10	10	10	10	10	10	10	10
	mass-dep. · empirical	11	11	11	11	11	11	11	11	11
	mass-dep. · tight	11	11	11	11	11	11	11	11	11
	long · empirical	9	9	9	9	9	9	16	16	16
	long · tight	11	11	11	11	11	11	11	11	16
Roh-Cha	short · empirical	0	0	0	0	0	0	0	0	0
	short · tight	0	0	0	0	0	0	0	0	0
	mass-dep. · empirical	0	0	0	0	0	0	0	0	0
	mass-dep. · tight	0	0	0	0	0	0	0	0	0
	long · empirical	0	0	0	0	0	0	0	0	0
	long · tight	0	0	0	0	0	0	0	0	0
Liu	short · empirical	0	0	0	0	0	0	0	0	0
	short · tight	0	0	0	0	0	0	0	0	0
	mass-dep. · empirical	0	0	0	0	0	0	0	0	0
	mass-dep. · tight	0	0	0	0	0	0	0	0	0
	long · empirical	0	0	0	0	0	0	0	0	0
	long · tight	0	0	0	0	0	0	0	0	0

Shift labour

Crew-team hours

- Colour:** the winning fleet; **number:** the share of blocks it couples (%)
 - Hatched:** winner kept in < 80% of day resamples (a near tie)
 - Red dot:** a decision that more search might flip
- 380 t ■ 425 t ■ 500 t ■ 550 t ■ 270 t + 1 heavy ■ 270 t + 2 heavy
 support < 80% ● weak search dependence

Example: the bottom six rows (Liu masses) are all dark blue, so 425 t is cheapest at every price, and the 0 means it never couples. Hover a cell for details.

Affine price curve. The row "heavy-tail · short handling · empirical due dates" has one known search miss: 270 t with two heavy units qualifies with one unit fewer and is cheaper than 550 t in 89 shift-labour decisions.

By mass profile: the smallest tier that carries (nearly) every block alone, or light units with a few heavy ones

MASS PROFILE	COVERING TIER	WINNERS (SHARE OF 1,080 DECISIONS)	COUPLES ≤ 1/10	COUPLES NONE	MEDIAN MARGIN
Jiang	500 t	425 t 42.1% · 270 t + 2 heavy 38.1% · 500 t 11.8% · 270 t + 1 heavy 5.6% · 300 t 2.2% · 380 t 0.2%	97.8%	11.8%	1.7%
Light-skewed	500 t	270 t + 2 heavy 59.2% · 500 t 24.2% · 270 t + 1 heavy 16.7%	100%	34.0%	1.6%
Uniform	500 t	500 t 99.3% · 270 t + 2 heavy 0.7%	99.3%	99.3%	2.0%
Heavy tail	—	270 t + 2 heavy 36.6% · 500 t 28.5% · 550 t 22.5% · 380 t 6.4% · 300 t 3.4% · 325 t 1.8% · 425 t 0.8%	47.6%	0%	1.2%
Roh–Cha	550 t	550 t 100%	100%	100%	30.1%
Liu	425 t	425 t 100%	100%	100%	3.3%

- In the five profiles with a covering tier, 99.4% of winners couple ≤ 1/10. In the Jiang sample only 1.3% of blocks exceed 425 t, so 425 t wins

■ **The heavy tail is the expected exception** (a menu gap): 9.1% of blocks exceed every tier, so even 550 and 500 t couple 9.1% and 10.8%; of the 145 shift-labour decisions whose winner couples > 1/10, 126 go to 500 t, which couples only the blocks above 500 t
- **Margins are often small:** the runner-up costs a median 2.4% more, under 1% in 18.3% of decisions. Yet outside the heavy tail every fleet within 1% of the least cost couples ≤ 1/10 in 98.8% of decisions (within 2%: 96.8%): the conclusions concern the kind of fleet, not one tonnage

■ Quotations: holding every rival and labour cost fixed, overturning the winner takes a capital-price increase of a median 11.1%, and under 5% in 1,123 decisions (17.3%), so purchases need comparable quotations

Sources: this study (paper §5.1, Table 5; Supplementary S9); sources of the mass profiles on “Parameter sources (1)”.

Try it: pick a condition and a price ratio and see which fleet is cheapest

Masses

Uniform

Handling

short handling

Due dates

empirical due dates

Shift labour

Crew-vehicle hours

Crew-team hours



Fleet	Composition	K*	Daily cost (labour-h)		vs cheapest	Coupled
500 t	5 × 500 t	5		421	★	0%
550 t	5 × 550 t	5		430	+2.0%	0%
270 t + 2 heavy	2 × 500 t + 4 × 270 t	6		474	+12.6%	11%
425 t	6 × 425 t	6		490	+16.4%	19%
270 t + 1 heavy	1 × 500 t + 6 × 270 t	7		543	+28.9%	33%
380 t	7 × 380 t	7		561	+33.2%	32%
300 t	8 × 300 t	8		619	+47.1%	51%
325 t	8 × 325 t	8		626	+48.7%	45%
250 t	9 × 250 t	9		681	+61.8%	63%
270 t	9 × 270 t	9		688	+63.3%	58%
200 t	12 × 200 t	12		888	+110.9%	76%

Cheapest: **500 t** (5 vehicles), 421 labour-hours a day; the runner-up, 550 t, costs 2.0% more. Labour is 76% of the cheapest fleet’s cost.

Capital


Labour

Daily cost $C = \theta \times P + L$ (labour-hours a day). P: fleet price in 270 t equivalents (affine curve); heavy units in mixes are priced at the smallest tier they need. Shift labour = 64 labour-hours per unit; crew measures count only hours actually worked.

Sources: counts and hours: this study; price curve fitted to 33 tender awards; 4 crew × 16 h: Park & Seo (2012), Roh & Cha (2011) (paper §3.2).

Question 2: coupled occupancy, alignment included, makes up most of the extra transporter-time; partner waiting is small

330 sized fleets outside the heavy-tail profile: transporter-hours a day above the covering fleet of the same condition



+1.6%

transporter-hours a day for each percentage point of blocks coupled
($r = 0.97$)

2.2×

the covering fleet's transporter-hours, for fleets that couple 60–100% of blocks

0.55%


waiting for the last team member: a median 0.55% of working time, never more than 2.2% for any fleet

+1.6% transporter-hours per point coupled ($r = 0.97$)

Partner waiting is small; the cost is that two or three vehicles and crews are held for the whole loaded trip, alignment included, and leave parallel service.

Block by block: a median 90% of a light fleet's extra transporter-time is coupled occupancy (repeated service plus alignment), 2% partner waiting

266 lighter fleets against the covering tier of their condition (outside the heavy tail, coupling > 1% of blocks): extra transporter-hours and their two parts



■ Coupled occupancy: repeated service plus alignment ■ Partner waiting — 1:1

89.8%

median share of the extra transporter-time that is coupled occupancy, $(|S| - 1) \cdot D + |S| \cdot \delta$ per block: repeated service plus alignment (10th–90th percentile 77.5–93.8%)


2.1%

median share spent waiting for partners (1.3–3.2%); early members could leave later, so this is an upper bound

- The rest is the difference in empty travel, which can be negative: points above the 1:1 line have less empty travel than the covering tier
- The same decomposition gives the covering fleet's occupancy per block used by the budget: about 43 min short, 63–85 mass-dependent, about 80 long loading

Calibrated on covering-fleet schedules, the fluid model rules out the lighter tier in 77% of comparisons with one error; scheduling settles the rest

Transporters needed beyond the workload rule K_p (share of K^* , mean with 95% interval)



$r = 0.97$

fluid occupancy factor against the simulated count ratio K_{lo}^* / K_{hi}^* (210 fleet pairs)

77% · 1


with a 10% margin (fluid cost ratio > 1.10) the screen rules out the lighter tier in 7,311 of 9,450 comparisons; the simulation finds it cheaper in only one, by 0.1%

- Of the 571 comparisons the lighter tier wins, 570 lie among the 2,139 left to scheduling; a 5% margin rules out 82.0% with 35 errors, no margin 92.3% with 274
- K_p is within one of K^* for 319 of 384 fleets (all 192 with empirical due dates); under tight due dates fleets coupling $> 60\%$ need 10.2% of K^* more, those coupling none 0.6% fewer: a coupled lift needs several vehicles free at once
- Priced on K_p : single-carry share 92.0% (against 91.0%) but the same winner in only 57.7%; 17.7% of the selections lack a qualifying schedule

The 77% is a within-sample check: occupancy comes from covering-fleet schedules, and a screen from inputs alone is untested.

Coupling-reliant fleets could win only if capital became far dearer than crews

Share of decisions won by fleets coupling $\leq 1/10$ as θ runs from 0.5 to 5,000 (conditions $\times$ 5 price curves $\times$ the labour measures named)



■ four measures, outside heavy tail ■ crew-team hours, outside heavy tail
■ four measures, all 36 conditions

With counts and schedules fixed, every fleet's cost is a straight line in θ , so the winner at any θ follows exactly from the lower envelope.

100% · 98.3%

outside the heavy tail: every decision up to $\theta \approx 8.7$, and 98.3% at $\theta = 30$


82%

even with capital alone ($\theta \rightarrow \infty$): sub-additive prices make a single carrier cheaper per block without crews (Proposition 1)

- In 82% of the 600 combinations no θ up to 5,000 makes a coupling-reliant fleet cheapest; 98 of the other 108 switch above $\theta = 30$ (median 87 labour-hours, about 1.4 crew-days)
- The few switches inside the calibrated range all occur under crew-team hours (dashed), which already assume one crew can drive a coupled set
- A crew-day now costs 2.1–12.5 times a 270 t unit's daily capital; driverless transporters could approach the switch

As heavy blocks become common, the cheapest fleet moves from light to light-plus-heavy to heavy, always coupling little

Cheapest fleet by the share φ of heavy blocks (350–540 t; the rest 100–300 t), affine price · shift labour; the same at all nine price ratios



- 300 t: couples every heavy block
- 270 t units + one or two 550 t
- 550 t: carries every block alone

- 300 t wins every decision at $\varphi = 0$ and 92–97% at $\varphi = 2\%$; from $\varphi = 5$ –10% light units with one or two heavy units take over
- 550 t first wins at $\varphi = 50\%$ with short and at $\varphi = 20\%$ with mass-dependent handling (all 45 shift-labour decisions); the true switch lies between these points and the next lower ones
- The budget explains both switches: 300 t stays within its budget up to $\varphi \approx 5\%$; the mixes lose to 550 t once the heavy workload exceeds two heavy units
- The simple formula ranks 300 vs 550 t correctly in 94.4% of 630 comparisons

2,502 / 2,520


decisions where the cheapest fleet couples at most one block in ten; the exceptions are the two-heavy mix at $\varphi = 35\%$ with mass-dependent handling under per-team crews (19.6%)

The winners' mean coupled share is at most 3.3% at every φ .

Richer mixes change the cheapest fleet but not the principle: 300 t plus 425 t units win all 8,000 cost scenarios for the Liu sample

Liu sample, short handling, 120-min cap; candidates = the 11 reference fleets plus extended mixes; 5 price curves $\times$ 400 values of $\theta \times 4$ labour measures = 8,000 cost scenarios.

Heavy 425 t units in the cheapest mix (share of scenarios)



44% of the Liu blocks exceed 300 t, so a mass-matched five-vehicle fleet would hold about 2.2 heavy units; the cheapest mixes couple at most 2.1% of blocks.

0.3-7.7% is how much more the best other candidate costs


1.3-7.7%

is how much more the homogeneous 425 t fleet, the winner among the reference types, now costs under shift labour

The number of heavy units follows the heavy workload; richer mixes change which fleet is cheapest, not the single-carry principle.

In 13 further conditions light units with matched heavy units win 95.3% of decisions, and the heavy count follows the heavy-share rule

Heavy units of the cheapest 300 t + covering-unit mix relative to the rule value K_H^* = $[K_{cov} \cdot h_H]$ (180 decisions per condition)



95.3%

of 2,340 decisions are won by 300 t units with covering heavy units, always coupling $\leq 1/10$; the best reference fleet costs a median 1.9% more (at most 10.2%). The exceptions, in the light-skewed profile, go to 270 t units with one or two heavy units, also light-heavy mixes


96.6%

have a heavy count within one unit of K_H^* ; the most frequent count equals K_H^* in 6 of the 13 conditions and is one off in the others

- Too few heavy units show the budget at work: uniform profile, tight due dates, one 500 t unit leaves 27% of blocks coupled and needs seven transporters, while three need five and couple 0.3%
- In four conditions 79 decisions (3.4%) pick $K_H^* + 2$, the largest tested; in three of them the next count up is the homogeneous covering fleet, already a candidate, so only the Liu sample with long loading (1.2%) could have a cheaper untested count

Service needs a lateness cap: with an on-time share alone, coupling fleets qualify with one unit fewer by deferring blocks 13–19 h

Fleets that need one unit fewer without a cap: the latest block of each over 30 days



6


fleets qualify with one unit fewer without a cap: their latest block is 13–19 h late

The deferred blocks are mostly coupled ones: of the blocks more than 4 h late, 40 of 59 (300 t) and 16 of 23 (325 t) under Jiang masses, and 24 of 26 (325 t) under Liu masses. In the mixes they are blocks queued for the scarce heavy units.

With an average target alone, coupled lifts can be pushed back indefinitely, which no production plan can absorb; tighter caps never lower a count.

With a 120-min cap, fleets that qualify by deferring are excluded, and least cost rises by only 0–1.85%

Shift labour; short handling, empirical due dates; 30 days per mass profile



Sources: this study (paper §5.4, Fig. 8b, c).

Masses	No cap	120-min cap	Cost
Jiang	5 × 300 t 12.5%	1 × 500 t + 4 × 270 t 6.3%	+0.2–1.7%
Uniform	2 × 500 t + 3 × 270 t 10.2%	2 × 500 t + 3 × 300 t 5.5%	+0.3–1.9%
Liu	2 × 425 t + 3 × 300 t 2.5%	2 × 425 t + 3 × 300 t 2.1%	0

- Under Jiang and uniform masses the cap changes the purchase: to mixes with heavy units that couple less
- Without a cap the winner couples $\leq 1/10$ in only 75% of their decisions; with any cap from 480 down to 30 min, in all of them
- Tightened to 30 min, cost rises by at most 4.7%; under Liu masses 2 × 425 t + 3 × 300 t is cheapest at every cap

The on-time target matters far less than the cap: at 90% or 98% the count barely moves

40 · 37

of 42 fleets keep their K^* at a 90% · 98% target; the others move by one transporter

93.3%

•

87.1%

of 1,260 decisions keep the 95% winner; outside the heavy tail every winner couples $\leq 1/10$ at both levels

Why: the integer step of one transporter

- At K^* the best schedules already keep at least 98% of blocks on time for 36 of the 42 fleets (a 37th falls one block short)
- One transporter fewer drops the pooled on-time share to 56–96%, and leaves at least one day without a schedule within the cap for 32 of the 41 fleets that can lose one

An average target cannot replace a limit on extreme lateness.

Flexible coupling adds little at the fleet level; only zero alignment time lets a light fleet win

Two coupling rules

Flexible (main rule): any minimal team. **Restricted**: mixes may couple only blocks that no transporter of the fleet can carry alone. Counts found under the restricted rule are pooled into the flexible results, since its schedules remain feasible.

exactly the same		87.0%
≤ 1% more		91.0%
≤ 2.5% more		94.9%

Change of least cost under the restricted rule (shift labour, 1,620 decisions); at most 9.4%. Zero for the uniform, heavy-tail, Roh–Cha and Liu profiles, largest for light-skewed (mean 1.4%). The mixes need flexible coupling (only 6 of 72 still qualify), but the decision hardly changes, because a single-carry fleet is nearly as cheap. Without pooling, search noise would make the restricted rule look cheaper in 8.6% of comparisons.

Alignment time

Jiang sample, short handling, empirical due dates, six fleet types: with no alignment time ($\delta = 0$) a 300 t fleet coupling 12.5% wins 77 of 180 cost scenarios; from $\delta = 2.5$ min it wins none, and the winner becomes 270 t with one heavy unit (coupling 6.3–6.9%), then with two heavy units from $\delta = 20$ min ($\leq 0.2\%$).


88.6% → 100%

single-carry winners in the one-factor runs (reduced candidate set): 88.6% at $\delta = 0$, 100% at $\delta = 20$ –40 min

Coupling pays only when it is almost free.

Capacity-removal estimate: least-cost redundancy is a second heavy unit or one more transporter, not coupling

On-time share over the 30 days after removing one unit of capacity (one fewer in a capacity class, start positions reassigned; 42 contending fleets in seven conditions)



◆ mix with one heavy unit (heavy capacity removed) ● covering homogeneous fleet ● other fleets

44-64% · 83-87%

outside the heavy tail: single-heavy mixes without their heavy capacity class · covering homogeneous fleets with one fewer unit

- Redundancy is estimated from the qualifying counts of the original and reduced compositions. At those counts, normal-operation schedules give 87.5% of 6,480 winners coupling $\leq 1/10$ (98.8% outside the heavy tail), at a median 8.8% above the nominal winner; outside the heavy tail no single-heavy mix wins (5.8% of nominal winners)
- The option to couple is worth about as much as nominally (90th-percentile upper bounds 6.4–6.7%); fleets with more units lost relatively less service in only 3 of 7 conditions

This is a capacity-composition estimate: start positions are reassigned rather than kept for the surviving vehicles, so it does not establish service after an arbitrary vehicle fails.

How reliable are the results? We looked for what could overturn them

① Other working days

Draw 30 days from the 30 **with replacement** and compare all fleets again; repeat 2,000 times (the bootstrap) and count how often the result holds.

1	2	3	4	5	6	7	8	9	10	11	12	13	14
15	16	17	18	19	20	21	22	23	24	25	26	27	28
29	30												
1	1	2	3	3	3	5	5	5	7	7	8	9	10
10	11	18	18	18	19	19	21	22	23	24	24	26	28
29	30												

Top: the original 30 days; bottom: one redraw,
yellow = days drawn more than once

② Distance to a flip

Counts come from search: with more search a loser might manage with one unit fewer and overtake. We compute how many on-time blocks it still lacks at one fewer and compare that with what two extra searches recover.

strong: short by > 92

medium: 59–92

weak: ≤ 58

92 = the most on-time blocks two extra searches recovered in any fleets (30 days); 58.9 = its 95th percentile

③ Disturb the inputs

Multiply every loading and unloading time by a random 0.5–1.5 ($\pm 50\%$) in 20 scenarios a day, to see whether fleets sized to the target still suffice, then re-size them to meet it under noise.

handling $\times U(0.5, 1.5) \cdot 20 + 20$ scenarios a day

Also checked: exact counts on small instances, 30 fresh days, six more searches a day at the boundary, counts from shallow search, the worst case of every coupling fleet losing a unit, unmeasured parameters, manufacturer speeds, capacity removal and a second yard. All 214,210 scheduling runs are kept and can be recomputed.

Question 4: the checks change which fleet is cheapest, not whether it carries alone

CHECK	WHAT CHANGES	SINGLE-CARRY SHARE (OUTSIDE THE HEAVY TAIL)	OTHER OUTCOME
As reported, shift labour	—	91.0% (100%)	1,620 decisions
As reported, all labour measures	four labour measures	90.8% (99.4%)	6,480 decisions
Search depth	counts after 0, 100 and 300 iterations; one search a day	90.2 · 88.8 · 89.3 · 90.8%	same winner 77.5 · 81.1 · 89.4 · 95.0%
Boundary re-search	six more searches a day at K* – 1 for 66 fleets and 10 controls	90.8%	K* fell for 3 fleets (no control); 242 decisions (3.7%) changed
Exact benchmark	80 single-batch slices; 70 fleets in 8 conditions	—	count equal to the exact count for 54 of 56 proven fleets, one higher for 2; certified winner in all 360 shift-labour decisions
Adversarial bound	all 258 coupling-reliant fleets given one transporter fewer	68.8% (80.9%)	worst case
Day resampling	2,000 bootstrap samples of the 30 days, stored schedules re-used	89.2% (> 99.9%)	K* unchanged in 89.0% of resamples on average (median 97%)
Fresh days	30 new days; 42 fleets in 7 conditions	98.6% (98.3%)	same winner in 86.5% of 1,260 decisions; K* within one for all 42
Slow speed and long handling	425 t, 500 t, mixes; handling ×1.5, ×2	100% of 3,060 resolved	540 decisions outside the specification for every candidate
Capacity removal	reduced compositions, starts reassigned	87.5% (98.8%)	same winner in 78.6%; least cost +8.8% (median)
Second case	Liu et al. (2022) data, sized from scratch	100%	425 t cheapest in 179 of 180 decisions

Search dependence is deliberately strict: a decision counts if any loser would be cheaper with one transporter fewer (5,033 of 6,480, 77.7%); this is not a share of wrong decisions. Graded at K* – 1: 3,513 strong, 298 medium, 1,222 weak. All 1,447 certain decisions have a winner coupling ≤ 1/10, against 88.1% of the search-dependent ones.

Sources: this study (paper §5.5, Table 6; Supplementary Table S1).

Against exact counts on small instances: the procedure finds the proven minimum for 54 of 56 fleets and is never below it

Each method's count against the proven minimum (single-batch slices, the 56 fleets with a proven minimum)

54 / 56


fleets whose count equals the proven minimum; 2 are one higher, none lower

360 / 360

shift-labour decisions pick a certified winner, although all are close

- One search run finds the fewest late blocks on 94.1% of 1,147 proven slice–count pairs, never more than one above
- Constructions alone or a greedy dispatcher: counts too high for 11 and 49 of 56 fleets (greedy: one more without a count); cheapest fleet changed in 37.5% and 25% of decisions, at up to 4.5% and 11.9% more cost


Good scheduling matters for fleet design: cruder dispatching overstates the fleet.



■ equals the proven minimum ■ one too many ■ two or more too many ■ no qualifying count

On 30 fresh days the kind of fleet holds; the named winner is the same in 86.5% of decisions

Share of each condition's 180 decisions whose least-cost fleet is the same as on the original days



39 / 42

fleets keep their K*; the rest stay within one transporter (two fell, one rose)

98.3%

of the 1,080 decisions outside the heavy tail have a winner coupling $\leq 1/10$ (98.6% of all 1,260)

- The largest change is in the heavy tail, where a 500 t fleet needed one transporter fewer and wins 84 of the 180 cost scenarios that 550 t or 270 t with two heavy units had won
- Holding the original counts, 41 of 42 fleets still have qualifying schedules on the fresh days; the exception is Roh–Cha 500 t with short handling: six transporters have no capped schedule on one day, seven qualify

Manufacturer speeds by model: 97% still couple $\leq 1/10$; only under Jiang masses with short handling does another heavy-unit mix win

Qualifying count K*: common speed → manufacturer speeds (yellow = changed; grey = not rerun, count kept)

Condition (empirical due dates)	380 t	425 t	500 t	550 t	270 t + 1 heavy	270 t + 2 heavy
Jiang · short	5 → 6	5 → 5	5 → 5	5 → 5	5 → 6	5 → 5
Jiang · mass-dep.	8 → 8	7 → 7	7 → 7	7 → 7	9 → 9	8 → 8
Roh–Cha · short	8 → 9	8*	6*	5 → 5	10*	8 → 9
Roh–Cha · mass-dep.	16 → 16	16*	12*	9 → 9	19*	17 → 18
Liu · short	6 → 7	5 → 5	5 → 5	5 → 5	8 → 9	6 → 7
Liu · mass-dep.	10 → 11	8 → 8	8 → 8	8 → 8	14*	12 → 13

380 and 425 t at 10/5 km/h, 500 t, 550 t and the mixes at 12/5 km/h. A faster vehicle never makes a block later, so common speeds are the optimistic end and mixes at the heavy speed the pessimistic end.

45 / 270

shift-labour decisions change their cheapest fleet, all with Jiang masses and short handling: the one-heavy mix needs one more vehicle and the two-heavy mix replaces it; keeping the old choice costs up to 18.6% more

178 / 1,080

decisions change their cheapest fleet at the pessimistic end (114 at the optimistic end)


1,054 / 1,080

decisions where the cheapest fleet still couples at most one block in ten (pessimistic end)

- Under Roh–Cha masses 550 t is cheapest in all 360 decisions; under Liu masses 425 t in 350 of 360 (among the eleven standard fleets)

Unmeasured parameters one at a time: the choice moves only between neighbouring fleets; only zero alignment time lifts coupling above one in ten

Settings whose cheapest fleet changed when one parameter changed (affine price · shift labour; right: the largest share the cheapest fleet coupled)



■ cheapest fleet changed ■ no fleet qualifies ■ couples > one in ten

Grey: no fleet qualified within four units above its starting count (all mass-dependent handling). With 425 t, 500 t and mixes added and scans up to 40 transporters, every resolved winner couples at most 6.2% of blocks (3,060 decisions); in the other 540 no candidate meets the specification at any count up to 40, so they lie outside the conclusions.

98.3%

of 10,260 decided decisions: the cheapest fleet couples $\leq 1/10$

12.5%


zero alignment time: 300 t becomes cheapest (Jiang masses, short handling), coupling 12.5%; from 2.5 min a mix with heavy units wins

- Every change is under Jiang masses, between neighbouring fleets; under Roh–Cha masses 550 t never changes
- With one crew per coupled set, 300 t also wins some Jiang decisions with longer handling or slower travel

Scope: 4 conditions (Jiang, Roh–Cha masses × short, mass-dependent handling) × 5 reference fleets (270, 300, 380, 550 t, 270 t + 2 heavy); alignment time was also run at 2.5 and 5 min with six fleets.

Fleets sized exactly to the target are lean: under $\pm 50\%$ handling noise most fall below 95%; re-sizing for noise changes the winner, but it still carries alone

On-time share under $\pm 50\%$ handling noise as units are added to K^* (one line per fleet; schedules chosen on 20 scenarios and evaluated on 20 others)



- At K^* : 89.0–97.4% with short and 79.8–93.1% with mass-dependent handling; only 15 of 66 fleets keep 95%
- Short handling: one extra unit suffices; mass-dependent: Jiang masses need one to three, and none of the 22 Roh-Cha and Liu fleets gets there within three

145 / 360

decided decisions change their winner: under Jiang masses, 270 t with two heavy units replaces one heavy unit (short handling) and 500 t replaces 425 t (mass-dependent)


≈ 0 blocks coupled by the winners under noise ($\leq 0.1\%$)

At the noise-robust counts, 42 of 44 fleets keep every block within the 120-min cap in every scenario; the other two exceed it on 0.8% of scenario-days.


A second yard from published data gives the same answer: 425 t is cheapest and couples nothing

Price ratio θ  12.4

Workload $\times 6$



Workload $\times 8$



A Shanghai cruise-ship yard's published week of transfers, distance table and five units of 250–420 t; demand $\times k = 1-8$, about $10 \cdot k$ moves a day, 30 days, 10-h shift. A reconstructed reference case, not a deployment.

- **Sized from scratch** ($k = 6, 8$): 425 t needs 5 and 6 units (270 t needs 9 and 14) and couples nothing; it is cheapest in 179 of the 180 decisions, by up to 10.6%
- **Adding to the existing five**: they suffice up to twice the demand; from $4\times$ without coupling, light additions never qualify while one or two 420 t units do (at $8\times$ with after-shift work)

The heavy units of the mixes are scheduled as 550 t vehicles and priced at 425 t. In the remaining decision, at six times the demand, 270 t units with two heavy units, coupling 7.6% of blocks, are 0.5% cheaper.

When does coupling pay? Four situations, and the limits of this study

Four situations where coupling pays or is needed

- **No tier can carry the block:** coupling is necessary (the heavy-tail profile)
- **Near-instant alignment:** at $\delta = 0$ a coupling 300 t fleet wins 43% of one condition's cost scenarios, at 2.5 min none
- **One crew drives a coupled set:** crew-team hours give the lowest single-carry share (86.9%) and, where a covering tier exists, the only switches inside the calibrated range
- **Capital far dearer than crews:** beyond about 1.4 crew-days per transporter-day, which remotely operated or automated transporters could approach

Under the conditions studied, least-cost purchasing favours single-carry capacity. Coupling is worth keeping for exceptional blocks, while the capacity-removal analysis favours a second heavy unit for redundancy.

Limitations

- Costs are capital and labour only; maintenance, energy and resale are omitted; capacity loss is tested with reduced compositions and reassigned start positions, so service after an arbitrary failure is unverified
- Labour uses four stylised measures, not a crew roster; heavy units of reference mixes are priced at the smallest tier they need (favouring the mixes)
- Stationary demand: seasonal peaks, rental and staged fleet expansion are outside the scope
- Static shortest paths, without traffic or conflicts between coupled sets; joint control, platform load distribution, road width, clearance, gradients and wheel loads are unmodelled (road width, clearance and wheel loads also constrain the heaviest single carriers, so their net effect is untested)
- Whole-day counts K^* are the smallest found, not certified minima (exact counts exist for single batches only)
- The conclusions exclude the slow-speed and long-handling cases where no candidate meets the specification

Status: all experiments are complete; the manuscript is being finalised

Done

Main study: 36 conditions × 11 fleet types Price and wage calibration
Search depth and boundary re-search Exact benchmark (single batches)
Day resampling and 30 fresh days Adversarial bound and evidence grading
Handling noise and robust counts Lateness cap and on-time targets 90% / 98%
Heavy-block share and extended mixes (16 conditions)
Restricted coupling and alignment time
Manufacturer speeds, one-factor sensitivity, slow speeds scanned to 40
Capacity-removal estimate Block-by-block decomposition Second yard

Next

- A final read of the revised manuscript against the result tables
- Archive the released code and result tables with a DOI
- Submit to Ocean Engineering

214,210

scheduling runs, about 2,570 core-hours, all reproducible

Five rules for a yard planning its transporter fleet

1

Make every frequent block carriable by one transporter

The least-cost fleet is the smallest tier that carries (nearly) every block alone, or light units with a few heavy ones; more capacity buys nothing. This qualifies the advice of Liu et al. (2022).

2

Use the coupling budget with calibrated occupancy

$h^* = \kappa \div (1 + 2\delta / \bar{w})$, about 2–18% under shift labour and $\alpha = 0.84$; with occupancy from covering-fleet schedules, exclude the lighter tier above a fluid cost ratio of 1.10.

3

Serve a thin heavy tail with a few heavy units, not coupling

Heavy units $[K_{\text{cov}} \cdot h_H]$, within one unit in 97% of the decisions tested.

4

Compare spare capacity with coupling for redundancy

Coupling is indispensable for blocks above every tier; in the capacity-removal estimate, least-cost redundancy is a second heavy unit or one more transporter. Failure protection needs checks that keep each surviving vehicle's capacity and position.

5

Specify a lateness cap; report near-optimal sets

An on-time share alone lets fleets defer their hardest blocks for hours; rankings are often within 1–2%, so report the near-optimal fleets and the decisions a better search could flip.

In one sentence: coupling is a remedy for exceptional blocks, not a purchasing strategy.

Thank you. Questions and comments are welcome.

References and data sources

1. Albert, B., Chaikovskaia, M., Gayon, J.-P., Quilliot, A. (2025). Sizing and scheduling for a fleet of reconfigurable mobile robots. *Int. J. Production Research* 63(18), 6758–6775. doi:10.1080/00207543.2025.2487565
2. Aziez, I., Côté, J.-F., Coelho, L.C. (2022). Fleet sizing and routing of healthcare automated guided vehicles. *Transportation Research Part E* 161, 102679. doi:10.1016/j.tre.2022.102679
3. Cao, Z., Zhang, S., Ceder, A. (2022). Novel coupling–decoupling strategy for scheduling autonomous public transport vehicles in overcrowded corridors. *Applied Mathematical Modelling* 106, 299–324. doi:10.1016/j.apm.2022.01.020
4. Chaikovskaia, M., Gayon, J.-P., Quilliot, A. (2026). Optimization of a fleet of reconfigurable robots. *Flexible Services and Manufacturing J.* 38(2), 470–492. doi:10.1007/s10696-025-09596-8
5. Drexl, M. (2012). Synchronization in vehicle routing. *Transportation Science* 46(3), 297–316. doi:10.1287/trsc.1110.0400
6. Efron, B., Tibshirani, R.J. (1993). *An Introduction to the Bootstrap*. Chapman & Hall.
7. Eskandari, H., Mahmoodi, E. (2016). A simulation-based multi-objective optimization study of the fleet sizing problem in the offshore industry. *Maritime Economics & Logistics* 18(4), 436–457. doi:10.1057/mel.2015.21
8. Fracapane, G., de Koster, R., Sgarbossa, F., Strandhagen, J.O. (2021). Planning and control of autonomous mobile robots for intralogistics: literature review and research agenda. *European J. Operational Research* 294(2), 405–426. doi:10.1016/j.ejor.2021.01.019
9. Gayon, J.-P., Lacomme, P., Oussama, A. (2026). Job-shop scheduling with cooperative transportation resources. *Int. Trans. in Operational Research*, Early View. doi:10.1111/itor.70145
10. Golden, B., Assad, A., Levy, L., Gheysens, F. (1984). The fleet size and mix vehicle routing problem. *Computers & Operations Research* 11(1), 49–66. doi:10.1016/0305-0548(84)90007-8
11. Guo, H., Wang, J., Sun, J., Mao, X. (2024). Multi-objective green vehicle scheduling in ship block transportation. *Scientific Reports* 14, 10796. doi:10.1038/s41598-024-61578-2
12. Hashemi Doulabi, H., Pesant, G., Rousseau, L.-M. (2020). Vehicle routing problems with synchronized visits and stochastic travel and service times: applications in healthcare. *Transportation Science* 54(4), 1053–1072. doi:10.1287/trsc.2019.0956
13. Hoff, A., Andersson, H., Christiansen, M., Hasle, G., Løkketangen, A. (2010). Fleet composition and routing. *Computers & Operations Research* 37(12), 2041–2061. doi:10.1016/j.cor.2010.03.015
14. Hu, Z.-H., Wei, C. (2018). Synchronizing vehicles for multi-vehicle and one-cargo transportation. *Computers & Industrial Engineering* 119, 36–49. doi:10.1016/j.cie.2018.03.023
15. Jiang, Z., Chen, Y., Li, X., Li, B. (2021). Multi-vehicle and one-cargo green transportation scheduling in shipbuilding. *Advanced Engineering Informatics* 49, 101306. doi:10.1016/j.aei.2021.101306
16. Joo, C.M., Kim, B.S. (2014). Block transportation scheduling under delivery restriction in shipyard. *Expert Systems with Applications* 41(6), 2851–2858. doi:10.1016/j.eswa.2013.10.020
17. Kim, B.S., Joo, C.M. (2012). Ant colony optimisation for block transportation scheduling with heterogeneous transporters. *Int. J. Production Research* 50(24), 7229–7241. doi:10.1080/00207543.2011.645078
18. Koç, Ç., Bektaş, T., Jabali, O., Laporte, G. (2016). Thirty years of heterogeneous vehicle routing. *European J. Operational Research* 249(1), 1–21. doi:10.1016/j.ejor.2015.07.020
19. Kweon, O., Kim, J., Kim, B., Kim, S.H. (2026). Pickup and delivery with time windows and two-dimensional loading for shipbuilding block transportation. *Computers & Operations Research* 194, 107573. doi:10.1016/j.cor.2026.107573
20. Liu, J., Yin, J., Khan, R.U. (2022). Scheduling management and optimization analysis of intermediate products transfer in a shipyard for cruise ships. *PLOS ONE* 17(3), e0265047. doi:10.1371/journal.pone.0265047
21. Nishi, T., Hiranaka, Y., Grossmann, I.E. (2011). A bilevel decomposition algorithm for simultaneous production scheduling and conflict-free routing for automated guided vehicles. *Computers & Operations Research* 38(5), 876–888. doi:10.1016/j.cor.2010.08.012
22. Pan, Z., Yu, Y. (2024). Optimizing heterogeneous multi-robot team composition for long-horizon construction tasks: time- and utilization-guided simulation. *Automation in Construction* 165, 105520. doi:10.1016/j.autcon.2024.105520
23. Park, C., Seo, J. (2012). A GRASP approach to transporter scheduling and routing at a shipyard. *Computers & Industrial Engineering* 63(2), 390–399. doi:10.1016/j.cie.2012.04.010
24. Pisinger, D., Ropke, S. (2007). A general heuristic for vehicle routing problems. *Computers & Operations Research* 34(8), 2403–2435. doi:10.1016/j.cor.2005.09.012
25. Roh, M.-I., Cha, J.-H. (2011). A block transportation scheduling system considering a minimisation of travel distance without loading of and interference between multiple transporters. *Int. J. Production Research* 49(11), 3231–3250. doi:10.1080/00207543.2010.484427
26. Ropke, S., Pisinger, D. (2006). An adaptive large neighborhood search heuristic for the pickup and delivery problem with time windows. *Transportation Science* 40(4), 455–472. doi:10.1287/trsc.1050.0135
27. Soares, R., Marques, A., Amorim, P., Parragh, S.N. (2024). Synchronisation in vehicle routing. *European J. Operational Research* 313(3), 817–840. doi:10.1016/j.ejor.2023.04.007
28. Sun, M., Liang, C., Chang, D. (2024). Dynamic scheduling using digital twin technology. *PLOS ONE* 19(2), e0297069. doi:10.1371/journal.pone.0297069
29. Vis, I.F.A. (2006). Survey of research in the design and control of automated guided vehicle systems. *European J. Operational Research* 170(3), 677–709. doi:10.1016/j.ejor.2004.09.020
30. Vis, I.F.A., de Koster, R., Savelsbergh, M.W.P. (2005). Minimum vehicle fleet size under time-window constraints at a container terminal. *Transportation Science* 39(2), 249–260. doi:10.1287/trsc.1030.0063
31. Voigt, S. (2025). A review and ranking of operators in adaptive large neighborhood search for vehicle routing problems. *European J. Operational Research* 322(2), 357–375. doi:10.1016/j.ejor.2024.05.033
32. Wang, C., Wang, K., Tao, J., Zhou, Y. (2020). Research on real-time optimal path planning model and algorithm for ship block transportation in shipyard. *J. Marine Science and Engineering* 8(12), 991. doi:10.3390/jmse8120991
33. Wang, N., Yin, J. (2026). Enhancing coastal industrial efficiency: a scheduling optimization framework for irregular assembly blocks in cruise ship manufacturing. *Frontiers in Marine Science* 13, 1843859. doi:10.3389/fmars.2026.1843859
34. Zhou, J., Wu, H., Weng, H., He, L., Li, W., Yang, T. (2025). Optimization of engineering vehicle scheduling in shipbuilding and repair yards based on the dual-cycle strategy. *Logistics* 9(4), 163. doi:10.3390/logistics9040163

Suzhou Dafang DCY270, DCY380, DCY500 transporter data sheets (szdf.com, archived 2026-09-24); TII Group self-propelled transporter brochure 0317-EN; Google OR-Tools CP-SAT solver.

33 public tender award notices (CSSC e-bidding platform, China Tendering Union Purchase platform, Sinochem e-bidding platform, 2022–2026), listed in Supplementary Table S4.

National Bureau of Statistics of China, 2025 average wages of urban employees (release of 15 May 2026); China Statistical Yearbook 2025, Table 4-12; Shanghai HRSS (2025), State Council 2019 No. 13, MOHRSS 2015 No. 71, MOHURD 2016 No. 74.